

## Generating Images with VAEs and GANs

- In this assignment, you will implement two models: a Variational Auto Encoder (VAE) and a Generative Adversarial Network (GAN), for the task of image generation.
- You will train these models on Fashion-MNIST, a synthetic dataset including the grayscale images of shoes, clothes, hats, etc.
- For each task, we only provide abstract instructions. You have the freedom to implement the model from scratch and customize helper functions for your convenience.
- We will primarily grade your work based on the quality of generated images by visualization.
- The GAN models usually need long-term training to produce high quality images. If you need GPUs to accelerate your experiment, please consider using Google's colab which provides free GPU runtime.

### ▼ Variational Auto Encoder (VAE)

- A VAE is an image generation model following the design of an auto-encoder, consisting of an encoder and a decoder trained by minimizing the reconstruction error. Unlike conventional Autoencoders, in which the bottleneck is constructed as reducing the capacity of the latent block, VAEs parameterize the latent space with a prior distribution, thereby providing a statistical interpretation of the reconstruction process.
- A VAE model optimizes the Evidence Lower Bound (ELBO) as follows:

$$P(x) \geq E_{q_{\theta}(z|x)}[q_{\phi}(x|z)] + D_{KL}(q_{\theta}(z|x) || p(z))$$

where  $q_{\theta}$ ,  $q_{\phi}$  are instantiated as encoder and decoder, respectively.  $p(z)$  is the prior distribution, which is usually chosen as Normal distribution  $\mathcal{N}(0, 1)$

- After training VAE, you are expected to generate two diversified samples (function provided), which should hold similar content as the input.

```
"""
Import PyTorch libraries.
"""

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable
import torch.optim as optim
import torchvision
from torchvision import datasets, transforms
from torchvision.utils import save_image
import matplotlib.pyplot as plt
import numpy as np
import random
```

```

"""
Enable CUDA if the GPU is available
"""
if torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

"""
Define the dataloader for the Fashion MNIST dataset.
"""
train_set = torchvision.datasets.FashionMNIST("./data", download=True, transform=
                                              transforms.Compose([transforms.ToTensor()]))
test_set = torchvision.datasets.FashionMNIST("./data", download=True, train=False, transform=
                                              transforms.Compose([transforms.ToTensor()]))
train_loader = torch.utils.data.DataLoader(train_set,
                                             batch_size=100)
test_loader = torch.utils.data.DataLoader(test_set,
                                           batch_size=100)

def output_label(label):
    output_mapping = {
        0: "T-shirt/Top",
        1: "Trouser",
        2: "Pullover",
        3: "Dress",
        4: "Coat",
        5: "Sandal",
        6: "Shirt",
        7: "Sneaker",
        8: "Bag",
        9: "Ankle Boot"
    }
    input = (label.item() if type(label) == torch.Tensor else label)
    return output_mapping[input]

image, label = next(iter(train_loader))
plt.imshow(image.squeeze(), cmap="gray")
print(output_label(label))
print(label)
print(image.shape)

```

```

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
26422272/? [00:04<00:00, 9060244.39it/s]

Extracting ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz to ./data/FashionMNIST/raw/

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
29696/? [00:00<00:00, 39850.35it/s]

Extracting ./data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw/

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
4422656/? [00:02<00:00, 3215173.46it/s]

Extracting ./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to ./data/FashionMNIST/raw/

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
6144/? [00:00<00:00, 155221.08it/s]


Extracting ./data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw/

```

Ankle Boot

9

```
torch.Size([1, 28, 28])
```



## ▼ Define VAE Model (TODO) [20 points]

- Define encoder, decoder, and reparameterization module
- Implement the forward pass of the VAE model

**[REDACTED]**

```
"""
```

A Convolutional Variational Autoencoder

```
"""
```

```

class VAE(nn.Module):
    def __init__(self):
        super().__init__()
        self.input_layer = nn.Linear(784, 392)
        self.linear_down_m = nn.Linear(392, 20)
        self.linear_down_v = nn.Linear(392, 20)
        self.linear_up1 = nn.Linear(20, 392)
        self.linear_up2 = nn.Linear(392, 784)
        self.sigmoid_activation = nn.Sigmoid()

    def encode(self, x):

```

```

        x = F.relu(self.input_layer(x))
        mean = self.linear_down_m(x)
        logvar = self.linear_down_v(x)
        return mean, logvar

    def reparameterize(self, mean, var):
        logvar = var.exp()
        sample = torch.FloatTensor(mean.size()).normal_().to(device)
        out = sample.mul(logvar).add(mean)
        return out

    def decode(self, sample):
        x = F.relu(self.linear_up1(sample))
        out = self.sigmoid_activation(self.linear_up2(x))
        return out

    def forward(self, x):
        mean, logvar = self.encode(x)
        sample = self.reparameterize(mean, logvar)
        decoded = self.decode(sample)
        return decoded, mean, logvar

```

## ▼ Complete the Training Loop (TODO) [15 points]

- Implement the complete training loop for the VAE model

```

"""
Initialize Hyperparameters
"""

batch_size = 128
learning_rate = 1e-3
num_epochs = 20

"""
Initialize the network and the Adam optimizer
"""

model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

"""
Training the network for a given number of epochs
"""

for epoch in range(num_epochs):
    model.train()
    train_loss1 = 0
    train_loss2 = 0
    for _, data in enumerate(train_loader):
        image, _ = data
        image = image.view(image.size(0), -1).to(device)
        optimizer.zero_grad()
        out, mean, logvar = model(image)

```

```

        out, mean, logvar = model(image)
        image_loss = nn.MSELoss(reduction='sum')(out, image)
        kloss = -0.5 * torch.sum(1 + logvar - mean.pow(2) - logvar.exp())
        loss = image_loss + kloss
        loss.backward()
        optimizer.step()
        train_loss1 += image_loss
        train_loss2 += kloss

    print('Epoch {}/{}'.format(epoch, num_epochs), (train_loss1+train_loss2))

```

```

❏ Epoch 0/20, loss 35.6633
Epoch 1/20, loss 24.6583
Epoch 2/20, loss 22.6981
Epoch 3/20, loss 21.8235
Epoch 4/20, loss 21.3217
Epoch 5/20, loss 20.9933
Epoch 6/20, loss 20.7431
Epoch 7/20, loss 20.5362
Epoch 8/20, loss 20.3614
Epoch 9/20, loss 20.2149
Epoch 10/20, loss 20.1003
Epoch 11/20, loss 19.9962
Epoch 12/20, loss 19.9068
Epoch 13/20, loss 19.8332
Epoch 14/20, loss 19.7591
Epoch 15/20, loss 19.6899
Epoch 16/20, loss 19.6402
Epoch 17/20, loss 19.5820
Epoch 18/20, loss 19.5392
Epoch 19/20, loss 19.4910

```

## ▼ Visualizing the output

- Display the VAE output by running multiple samples of the latent space.
- You can modify this function accordingly to fit your implementation

```

"""

```

The following part takes two images from test loader to feed into the VAE.  
Both the original image and generated image(s) from the distribution are shown.

```

"""

```

```

import matplotlib.pyplot as plt
import numpy as np

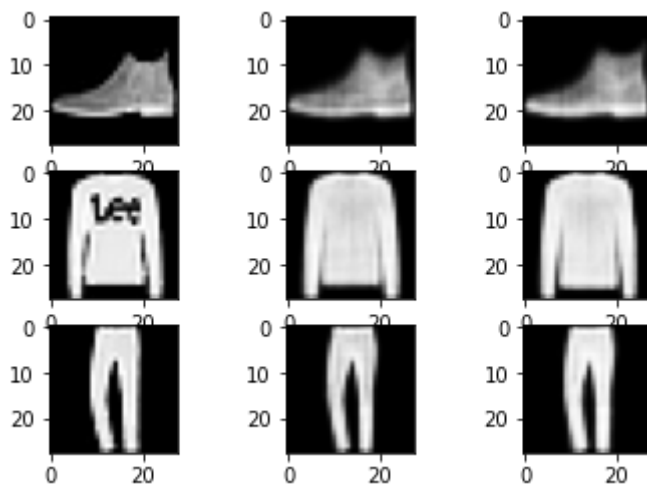
model.eval()
with torch.no_grad():
    imgs, _ = list(test_loader)[0]
    imgs = imgs.to(device)
    new_img = imgs.view(imgs.size(0), -1)

```

```

fig, ax = plt.subplots(3,3)
ax = ax.reshape(-1)
for i in range(3):
    for j in range(3):
        if j == 0:
            # input image
            img_i = np.transpose(imgs[i].cpu().numpy(), [1,2,0])
            ax[i*3+j].imshow(np.squeeze(img_i), cmap = 'gray')
        else:
            # vae generation results
            out, _, _ = model(new_img)
            out = out.cpu().numpy()[i].reshape(28,28)
            ax[i*3+j].imshow(out, cmap = 'gray')
plt.show()

```



## Generative Adversarial Network (GAN)

- A GAN is a generation model trained to convert the samples from prior distribution  $z$  to the target domains  $x$ , e.g., images or text, in an unsupervised fashion.
- A GAN consists of a generator ( $G$ ) and discriminator ( $D$ ) model, where  $G$  is trained to produce realistic samples of the target domain, while the  $D$  learns to identify the samples from generator and the real domain, and serve as the supervision to optimize the  $G$ .
- GAN training typically follows two iterative steps:

1.  $\max_D \log(D(x)) + \log(1 - D(G(z)))$
2.  $\max_G \log(D(G(z)))$

### ▼ Define GAN Model (TODO) [20 points]

- Define the generator and discriminator for the GAN model

```
start_size = 64
```

```

image_size = 784
int_size = 256
batch_size = 100

"""
Define the generator
"""
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()

        self.gen_seq = nn.Sequential(
            nn.Linear(start_size, int_size),
            nn.ReLU(),
            nn.Linear(int_size, int_size),
            nn.ReLU(),
            nn.Linear(int_size, image_size),
            nn.Tanh()
        )
    def forward(self, input):
        output = self.gen_seq(input)
        return output

"""
Define the discriminator
"""
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.disc_seq = nn.Sequential(
            nn.Linear(image_size, int_size),
            nn.LeakyReLU(0.2),
            nn.Linear(int_size, int_size),
            nn.LeakyReLU(0.2),
            nn.Linear(int_size, 1),
            nn.Sigmoid()
        )
    def forward(self, input):
        output = self.disc_seq(input)
        return output

```

#### ▼ Complete the Training Loop (TODO) [15 points]

- Implement the complete training loop for the GAN model
- You can monitor the loss curve between generator and discriminator to tune the hyperparamters

- The loss of generator and discriminator converge toward some non-zero values at the end.

```

"""
Initialize Hyperparameters
"""
learning_rate = 0.0002
num_epochs = 300
"""
Initialize the network and the Adam optimizer
"""
generator = Generator().to(device)
discriminator = Discriminator().to(device)
loss_func = nn.BCELoss()
disc_optimizer = torch.optim.Adam(discriminator.parameters(), lr=learning_rate)
gen_optimizer = torch.optim.Adam(generator.parameters(), lr=learning_rate)

disc_losses = []
gen_losses = []
DofXs = []
GofDofXs = []

for epoch in range(num_epochs):
    for i, (images, label) in enumerate(train_loader):
        images = images.reshape(batch_size, -1).to(device)

        #discriminator
        rlabels = torch.ones(batch_size, 1).to(device)
        flabels = torch.zeros(batch_size, 1).to(device)
        out = discriminator(images)
        discriminator_rloss = loss_func(out, rlabels)
        DofX = out

        z = torch.randn(batch_size, start_size).to(device)
        fakes = generator(z)
        out = discriminator(fakes)
        discriminator_floss = loss_func(out, flabels)
        GofDofX = out

        discriminator_loss = discriminator_rloss + discriminator_floss

        disc_optimizer.zero_grad()
        gen_optimizer.zero_grad()
        discriminator_loss.backward()
        disc_optimizer.step()

    #generator
    z = torch.randn(batch_size, start_size).to(device)
    fakes = generator(z)
    zlabels = torch.ones(batch_size, 1).to(device)
    generator_loss = loss_func(discriminator(fakes), zlabels)

```



```

disc_optimizer.zero_grad()
gen_optimizer.zero_grad()
generator_loss.backward()
gen_optimizer.step()

if (i+1) % len(train_loader) == 0:
    disc_losses.append(discriminator_loss.item())
    gen_losses.append(generator_loss.item())
    DofXs.append(DofX.mean().item())
    GofDofXs.append(GofDofX.mean().item())
    print('Epoch {}/{}'.format(epoch, num_epochs), disc_loss: {:.4f}, gen_loss: {:.4f}, D(x): {:.4f}, D(G(x)): {:.4f})

```

```

Epoch 240/300, disc_loss: 0.9709, gen_loss: 1.5238, D(x): 0.6672, D(G(x)): 0.2651
Epoch 241/300, disc_loss: 0.9592, gen_loss: 1.7260, D(x): 0.6908, D(G(x)): 0.3011
Epoch 242/300, disc_loss: 0.8626, gen_loss: 1.7123, D(x): 0.7225, D(G(x)): 0.2731
Epoch 243/300, disc_loss: 0.7730, gen_loss: 2.1976, D(x): 0.7064, D(G(x)): 0.2121
Epoch 244/300, disc_loss: 1.0898, gen_loss: 2.1009, D(x): 0.6341, D(G(x)): 0.2481
Epoch 245/300, disc_loss: 0.8978, gen_loss: 1.8393, D(x): 0.6734, D(G(x)): 0.1921
Epoch 246/300, disc_loss: 0.9226, gen_loss: 1.8197, D(x): 0.7173, D(G(x)): 0.3041
Epoch 247/300, disc_loss: 0.9041, gen_loss: 1.6426, D(x): 0.6764, D(G(x)): 0.2521
Epoch 248/300, disc_loss: 1.1963, gen_loss: 2.0043, D(x): 0.6264, D(G(x)): 0.2741
Epoch 249/300, disc_loss: 0.9203, gen_loss: 2.1373, D(x): 0.6891, D(G(x)): 0.2501
Epoch 250/300, disc_loss: 0.8779, gen_loss: 1.6488, D(x): 0.7002, D(G(x)): 0.2791
Epoch 251/300, disc_loss: 0.9717, gen_loss: 1.9046, D(x): 0.6728, D(G(x)): 0.2681
Epoch 252/300, disc_loss: 0.8281, gen_loss: 1.9360, D(x): 0.7560, D(G(x)): 0.3021
Epoch 253/300, disc_loss: 0.7616, gen_loss: 1.9395, D(x): 0.7432, D(G(x)): 0.2421
Epoch 254/300, disc_loss: 1.1080, gen_loss: 1.4316, D(x): 0.6817, D(G(x)): 0.3241
Epoch 255/300, disc_loss: 0.7856, gen_loss: 1.9725, D(x): 0.6817, D(G(x)): 0.1641
Epoch 256/300, disc_loss: 1.0327, gen_loss: 1.6056, D(x): 0.6816, D(G(x)): 0.3081
Epoch 257/300, disc_loss: 0.7766, gen_loss: 1.5415, D(x): 0.7382, D(G(x)): 0.2831
Epoch 258/300, disc_loss: 0.7791, gen_loss: 1.7538, D(x): 0.7779, D(G(x)): 0.2911
Epoch 259/300, disc_loss: 1.0993, gen_loss: 1.6785, D(x): 0.7320, D(G(x)): 0.3541
Epoch 260/300, disc_loss: 0.7157, gen_loss: 2.2125, D(x): 0.7440, D(G(x)): 0.2211
Epoch 261/300, disc_loss: 0.9472, gen_loss: 2.3232, D(x): 0.6360, D(G(x)): 0.2221
Epoch 262/300, disc_loss: 0.6461, gen_loss: 1.9910, D(x): 0.7452, D(G(x)): 0.1971
Epoch 263/300, disc_loss: 0.8125, gen_loss: 1.7770, D(x): 0.7097, D(G(x)): 0.2261
Epoch 264/300, disc_loss: 0.9019, gen_loss: 1.8847, D(x): 0.6926, D(G(x)): 0.2331
Epoch 265/300, disc_loss: 0.8331, gen_loss: 1.8891, D(x): 0.7316, D(G(x)): 0.2521
Epoch 266/300, disc_loss: 0.7838, gen_loss: 1.7998, D(x): 0.7455, D(G(x)): 0.2591
Epoch 267/300, disc_loss: 0.9543, gen_loss: 1.7060, D(x): 0.6724, D(G(x)): 0.2491
Epoch 268/300, disc_loss: 0.8830, gen_loss: 2.1059, D(x): 0.6941, D(G(x)): 0.2451
Epoch 269/300, disc_loss: 0.8783, gen_loss: 1.8287, D(x): 0.6860, D(G(x)): 0.2301
Epoch 270/300, disc_loss: 1.0270, gen_loss: 1.3826, D(x): 0.7199, D(G(x)): 0.3551
Epoch 271/300, disc_loss: 0.8477, gen_loss: 2.0885, D(x): 0.7403, D(G(x)): 0.2801
Epoch 272/300, disc_loss: 0.7030, gen_loss: 1.8038, D(x): 0.7415, D(G(x)): 0.2191
Epoch 273/300, disc_loss: 0.9698, gen_loss: 1.8164, D(x): 0.6632, D(G(x)): 0.2531
Epoch 274/300, disc_loss: 0.8587, gen_loss: 1.7720, D(x): 0.7488, D(G(x)): 0.3301
Epoch 275/300, disc_loss: 0.8598, gen_loss: 1.6185, D(x): 0.6891, D(G(x)): 0.2241
Epoch 276/300, disc_loss: 0.9355, gen_loss: 1.8444, D(x): 0.6949, D(G(x)): 0.2681
Epoch 277/300, disc_loss: 1.0768, gen_loss: 1.8741, D(x): 0.7328, D(G(x)): 0.3871
Epoch 278/300, disc_loss: 0.8365, gen_loss: 1.7646, D(x): 0.6982, D(G(x)): 0.2601
Epoch 279/300, disc_loss: 0.8664, gen_loss: 1.6122, D(x): 0.6776, D(G(x)): 0.2261
Epoch 280/300, disc_loss: 0.9077, gen_loss: 1.5876, D(x): 0.7140, D(G(x)): 0.2941
Epoch 281/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 282/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 283/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 284/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 285/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 286/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 287/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 288/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 289/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 290/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 291/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 292/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 293/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 294/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 295/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 296/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 297/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 298/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761
Epoch 299/300, disc_loss: 0.8640, gen_loss: 1.8255, D(x): 0.7347, D(G(x)): 0.2761

```

```

Epoch 281/300, disc_loss: 0.8642, gen_loss: 1.8355, D(x): 0.7247, D(G(x)): 0.2761
Epoch 282/300, disc_loss: 0.8788, gen_loss: 1.6858, D(x): 0.7614, D(G(x)): 0.3287
Epoch 283/300, disc_loss: 0.9690, gen_loss: 1.7742, D(x): 0.7249, D(G(x)): 0.3437
Epoch 284/300, disc_loss: 0.8060, gen_loss: 2.0578, D(x): 0.7346, D(G(x)): 0.2576
Epoch 285/300, disc_loss: 1.0217, gen_loss: 1.5658, D(x): 0.7010, D(G(x)): 0.3421
Epoch 286/300, disc_loss: 0.9670, gen_loss: 1.4514, D(x): 0.7311, D(G(x)): 0.3440
Epoch 287/300, disc_loss: 0.9032, gen_loss: 1.6142, D(x): 0.7897, D(G(x)): 0.3801
Epoch 288/300, disc_loss: 0.9412, gen_loss: 1.7596, D(x): 0.7114, D(G(x)): 0.3010
Epoch 289/300, disc_loss: 0.7955, gen_loss: 1.7024, D(x): 0.8027, D(G(x)): 0.3291
Epoch 290/300, disc_loss: 0.6479, gen_loss: 1.8811, D(x): 0.7693, D(G(x)): 0.2258
Epoch 291/300, disc_loss: 0.8711, gen_loss: 1.7604, D(x): 0.7001, D(G(x)): 0.2851
Epoch 292/300, disc_loss: 0.7225, gen_loss: 1.9267, D(x): 0.7543, D(G(x)): 0.2561
Epoch 293/300, disc_loss: 0.9898, gen_loss: 1.4098, D(x): 0.7146, D(G(x)): 0.3311
Epoch 294/300, disc_loss: 0.8634, gen_loss: 1.6264, D(x): 0.7213, D(G(x)): 0.2931
Epoch 295/300, disc_loss: 1.0210, gen_loss: 1.6315, D(x): 0.7244, D(G(x)): 0.3501
Epoch 296/300, disc_loss: 0.9259, gen_loss: 1.9647, D(x): 0.7298, D(G(x)): 0.3241
Epoch 297/300, disc_loss: 0.9731, gen_loss: 1.4594, D(x): 0.6752, D(G(x)): 0.2741

```


## ▼ Visualizing the output

- Display the GAN output by running multiple samples of the latent space.
- You can modify this function accordingly to fit your implementation

```

import matplotlib.pyplot as plt
fig, ax = plt.subplots(3,3)
ax = ax.reshape(-1)
for i in range(9):
    z = torch.randn(1, latent_size).to(device)
    fake = generator(z)
    fake = fake.reshape((-1, 28,28)).squeeze(0).detach()
    ax[i].imshow(fake.cpu().data.numpy(), cmap = 'gray')
plt.show()

```



---

✓ 3s completed at 12:18 AM

